



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Neural Networks

Leonid Mill, Tobias Würfl, Katharina Breining, Nishant Ravikumar, Sebastian Gündel, Mathis Hoffmann, Florian Thamm, Felix Denzinger
Pattern Recognition Lab, Friedrich-Alexander University of Erlangen-Nürnberg
October 25, 2018



Flexibility vs. abstraction

Low level

High level



- Linear Algebra operations
- Bare metal





- Compiles graphs of Tensor operations
- High flexibility



- Stacks together elementary layers
- Reduced flexibility

Artificial Neural Networks



$$y = f\left(\sum_N w_i x_i\right)$$

Terminology

- We will call $\frac{\partial L}{\partial \hat{\mathbf{y}}}$ the **error E** in the exercises
- "Layer" it is now a technical term. Layers must not be present in graphical depictions. E.g. activation functions become "layers"



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Neural Network



Neural Network

In layer-oriented frameworks we typically have a neural network object which

Neural Network

In layer-oriented frameworks we typically have a neural network object which

- is responsible for holding a **graph of layers**
 - we allow only extremely simple graphs
 - with a list of layers
 - and only one data source
 - and one loss function

Neural Network

In layer-oriented frameworks we typically have a neural network object which

- is responsible for holding a **graph of layers**
 - we allow only extremely simple graphs
 - with a list of layers
 - and only one data source
 - and one loss function
- is responsible to hold **access to data**

Neural Network

In layer-oriented frameworks we typically have a neural network object which

- is responsible for holding a **graph of layers**
 - we allow only extremely simple graphs
 - with a list of layers
 - and only one data source
 - and one loss function
- is responsible to hold **access to data**
- has **no explicit knowledge** about the graph of layers it contains

Neural Network

In layer-oriented frameworks we typically have a neural network object which

- is responsible for holding a **graph of layers**
 - we allow only extremely simple graphs
 - with a list of layers
 - and only one data source
 - and one loss function
- is responsible to hold **access to data**
- has **no explicit knowledge** about the graph of layers it contains
- **recursively calls forward** on its layers passing the input-data
- **recursively calls backward** on its layers passing the error

Neural Network

In layer-oriented frameworks we typically have a neural network object which

- is responsible for holding a **graph of layers**
 - we allow only extremely simple graphs
 - with a list of layers
 - and only one data source
 - and one loss function
- is responsible to hold **access to data**
- has **no explicit knowledge** about the graph of layers it contains
- **recursively calls forward** on its layers passing the input-data
- **recursively calls backward** on its layers passing the error
- in our case it stores the loss over iterations, while in other frameworks this is commonly separated into an optimizer class




FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING


Fully Connected Layer



Forward




Forward



$$\begin{pmatrix} w_1 \\ \vdots \\ w_n \end{pmatrix}^T \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} + w_{n+1} = \hat{y}$$

$$\mathbf{w}^T \mathbf{x} = \hat{y}$$


Forward




$$\begin{pmatrix} w_1 \\ \vdots \\ w_n \\ w_{n+1} \end{pmatrix}^T \begin{pmatrix} x_1 \\ \vdots \\ x_n \\ 1 \end{pmatrix} = \hat{y}$$

$$\mathbf{w}^T \mathbf{x} = \hat{y}$$

Forward




Forward



$$\begin{pmatrix} w_{1,1} & \dots & w_{1,m} \\ \vdots & \ddots & \vdots \\ w_{n,1} & \dots & w_{n,m} \\ w_{n+1,1} & \dots & w_{n+1,m} \end{pmatrix}^T \begin{pmatrix} x_1 \\ \vdots \\ x_n \\ 1 \end{pmatrix} = \begin{pmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_m \end{pmatrix}$$

$$\mathbf{W}\mathbf{x} = \hat{\mathbf{y}}$$

Forward



$$\begin{pmatrix} w_{1,1} & \dots & w_{1,m} \\ \vdots & \ddots & \vdots \\ w_{n,1} & \dots & w_{n,m} \\ w_{n+1,1} & \dots & w_{n+1,m} \end{pmatrix}^T \begin{pmatrix} x_{1,1} & \dots & x_{1,b} \\ \vdots & \ddots & \vdots \\ x_{n,1} & \dots & x_{n,b} \\ 1 & \dots & 1 \end{pmatrix}$$

$$\mathbf{WX} = \hat{\mathbf{Y}} \quad (1)$$

Backward

- Return gradient with respect to **X**:

Backward

- Return gradient with respect to **X**:

$$\mathbf{E}_{n-1} = \mathbf{W}^T \mathbf{E}_n \quad (2)$$

- \mathbf{E}_n : **error_tensor** passed downward

Backward

- Return gradient with respect to **X**:

$$\mathbf{E}_{n-1} = \mathbf{W}^T \mathbf{E}_n \quad (2)$$

- Update **W** using gradient with respect to **W**:

- **E_n**: **error_tensor** passed downward

Backward

- Return gradient with respect to **X**:

$$\mathbf{E}_{n-1} = \mathbf{W}^T \mathbf{E}_n \quad (2)$$

- Update **W** using gradient with respect to **W**:


$$\mathbf{W}^{t+1} = \mathbf{W}^t - \delta \cdot \mathbf{E}_n \mathbf{X}^T \quad (3)$$

Note: Dynamic programming part of Backpropagation


- E_n**: **error_tensor** passed downward
- δ : learning rate **delta** individual to this layer

Memory Layout

- Numpy uses C ordering by default
- Wrong ordering will cause strided data access
- We want the batch size to be the outermost loop
→ We have to adjust our formulas for the implementation



Forward - Our Memory Layout



$$\begin{pmatrix} x_{1,1} & \dots & x_{1,b} \\ \vdots & \ddots & \vdots \\ x_{n,1} & \dots & x_{n,b} \\ 1 & \dots & 1 \end{pmatrix}^T \begin{pmatrix} w_{1,1} & \dots & w_{1,m} \\ \vdots & \ddots & \vdots \\ w_{n,1} & \dots & w_{n,m} \\ w_{n+1,1} & \dots & w_{n+1,m} \end{pmatrix}$$

$$\mathbf{x}'\mathbf{w}' = \hat{\mathbf{y}}' \quad (4)$$

with

$$\mathbf{x}' = \mathbf{x}^T, \mathbf{w}' = \mathbf{w}^T, \hat{\mathbf{y}}' = \hat{\mathbf{y}}^T \quad (5)$$

$$\hat{\mathbf{y}}^T = (\mathbf{w}\mathbf{x})^T = \mathbf{x}^T\mathbf{w}^T \quad (6)$$

Backward - Our Memory Layout

- Return gradient with respect to **X**:

$$\mathbf{E}'_{n-1} = \mathbf{E}'_n \mathbf{W}'^T \quad (7)$$

- Update **W'** using gradient with respect to **W'**:

$$\mathbf{W}'^{t+1} = \mathbf{W}'^t - \delta \cdot \mathbf{X}'^T \mathbf{E}'_n \quad (8)$$

Note: Dynamic programming part of Backpropagation

- E'**_n: **error_tensor** passed downward
- δ : learning rate **delta** individual to this layer




FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

ReLU Activation Function



Forward



Backward

ReLU is not continuously differentiable!

Backward

ReLU is not continuously differentiable!

$$e_{n-1} = \begin{cases} 0 & \text{if } x \leq 0 \\ e_n & \text{else} \end{cases} \quad (9)$$

Note: DP part of Backpropagation yet again

Backward

ReLU is not continuously differentiable!

$$e_{n-1} = \begin{cases} 0 & \text{if } x \leq 0 \\ e_n & \text{else} \end{cases} \quad (9)$$

Note: DP part of Backpropagation yet again

- The scalar e is because activation functions operate elementwise on \mathbf{E}

Backward

ReLU is not continuously differentiable!

$$e_{n-1} = \begin{cases} 0 & \text{if } x \leq 0 \\ e_n & \text{else} \end{cases} \quad (9)$$

Note: DP part of Backpropagation yet again

- The scalar e is because activation functions operate elementwise on \mathbf{E}

- If you wonder about e_n instead of 1 consider that this is $\underbrace{\frac{\partial L}{\partial \hat{\mathbf{y}}}}_{\mathbf{E}} \cdot \underbrace{\frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{x}}}_{\text{ReLU}}$



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

"SoftMax Loss" Function



Forward

Labels as N -dimensional **one hot** vector \mathbf{y} :

$$\begin{pmatrix} \vdots \\ 1 \\ \vdots \end{pmatrix}$$

Forward

Labels as N -dimensional **one hot** vector \mathbf{y} : $\begin{pmatrix} \vdots \\ 1 \\ \vdots \end{pmatrix}$

- Activation(Prediction) $\hat{\mathbf{y}}$ for every element of the batch of size B :

$$\hat{y}_k = \frac{\exp(x_k)}{\sum_{j=1}^N \exp(x_j)} \quad (10)$$

Forward

Labels as N -dimensional **one hot** vector \mathbf{y} : $\begin{pmatrix} \vdots \\ 1 \\ \vdots \end{pmatrix}$

- Activation(Prediction) $\hat{\mathbf{y}}$ for every element of the batch of size B :

$$\hat{y}_k = \frac{\exp(x_k)}{\sum_{j=1}^N \exp(x_j)} \quad (10)$$

- Loss:

$$loss = \sum_{b=1}^B -\log \hat{y}_k \textbf{ where } y_k = 1 \quad (11)$$

Numeric

- If $x_k > 0 \rightarrow e^{x_k}$ might become very large
- To increase numerical stability x_k can be shifted
- $\tilde{x}_k = x_k - \max(\mathbf{x})$
- This leaves the scores unchanged!

Backward

For every element of the batch:

$$e_k = \begin{cases} \hat{y}_k - 1 & \text{where } y_k = 1 \\ \hat{y}_k & \text{else} \end{cases} \quad (12)$$

Backward

For every element of the batch:

$$e_k = \begin{cases} \hat{y}_k - 1 & \text{where } y_k = 1 \\ \hat{y}_k & \text{else} \end{cases} \quad (12)$$

- We simply decrease the activation of the correct class
- And increase all others

Backward

For every element of the batch:

$$e_k = \begin{cases} \hat{y}_k - 1 & \text{where } y_k = 1 \\ \hat{y}_k & \text{else} \end{cases} \quad (12)$$

- We simply decrease the activation of the correct class
- And increase all others
- Notice that this does **not** depend on an error **E**
- Because it's the starting point of the recursive computation of gradients



Thanks for listening.
Any questions?